

IT350 DA Assignment 5

MNIST Dataset

Submitted by: Harsh Agarwal (181IT117)

```
In [1]: # Import Libraries
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout
import matplotlib.pyplot as plt
import numpy as np
import matplotlib.pyplot as plt
np.random.seed(2017)
```

```
In [2]: # the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11495376/11490434 [=====] - 8s 0us/step
```

```
In [3]: #Flatten image
x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
#Normalize pixel data
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

60000 train samples
10000 test samples
```

```
In [4]: # convert class vectors to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
In [5]: batch_size = 128
epochs = 20
```

```
In [6]: #Initialise model
model = Sequential()
#Add input and 1st hidden layer
model.add(Dense(512, activation='sigmoid', input_shape=(784,)))
#Add 2nd hidden layer
model.add(Dense(512, activation='sigmoid'))
#Add output layer
model.add(Dense(num_classes, activation='sigmoid'))
```

```
In [7]: #Print model information
model.summary()

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
dense (Dense)                (None, 512)                 401920

dense_1 (Dense)              (None, 512)                 262656

dense_2 (Dense)              (None, 10)                  5130
=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0
```

```
In [8]: #Compiling model: Using categorical crossentropy loss function and sgd optimiser
model.compile(loss='categorical_crossentropy',
              optimizer=keras.optimizers.SGD(lr=0.1, momentum=0.0, nesterov=False),
              metrics=['accuracy'])
```

```
In [9]: #Training model
history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))

#Evaluating model
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

Epoch 1/20
469/469 [=====] - 5s 5ms/step - loss: 2.0702 - accuracy: 0.2696 - val_loss: 0.8628 - val_accuracy: 0.7363
Epoch 2/20
469/469 [=====] - 2s 4ms/step - loss: 0.6864 - accuracy: 0.8133 - val_loss: 0.4591 - val_accuracy: 0.8733
Epoch 3/20
469/469 [=====] - 2s 4ms/step - loss: 0.4458 - accuracy: 0.8734 - val_loss: 0.3792 - val_accuracy: 0.8905
Epoch 4/20
469/469 [=====] - 2s 4ms/step - loss: 0.3847 - accuracy: 0.8892 - val_loss: 0.3415 - val_accuracy: 0.9010
Epoch 5/20
469/469 [=====] - 2s 4ms/step - loss: 0.3508 - accuracy: 0.8978 - val_loss: 0.3167 - val_accuracy: 0.9060
Epoch 6/20
469/469 [=====] - 2s 4ms/step - loss: 0.3288 - accuracy: 0.9035 - val_loss: 0.3075 - val_accuracy: 0.9105
Epoch 7/20
469/469 [=====] - 2s 4ms/step - loss: 0.3124 - accuracy: 0.9080 - val_loss: 0.2991 - val_accuracy: 0.9143
Epoch 8/20
469/469 [=====] - 2s 4ms/step - loss: 0.3109 - accuracy: 0.9099 - val_loss: 0.2927 - val_accuracy: 0.9150
Epoch 9/20
469/469 [=====] - 2s 4ms/step - loss: 0.3025 - accuracy: 0.9119 - val_loss: 0.2886 - val_accuracy: 0.9140
Epoch 10/20
469/469 [=====] - 2s 4ms/step - loss: 0.2856 - accuracy: 0.9155 - val_loss: 0.2794 - val_accuracy: 0.9201
Epoch 11/20
469/469 [=====] - 2s 4ms/step - loss: 0.2796 - accuracy: 0.9196 - val_loss: 0.2903 - val_accuracy: 0.9144
Epoch 12/20
469/469 [=====] - 2s 4ms/step - loss: 0.2800 - accuracy: 0.9203 - val_loss: 0.2689 - val_accuracy: 0.9221
Epoch 13/20
469/469 [=====] - 2s 4ms/step - loss: 0.2753 - accuracy: 0.9193 - val_loss: 0.2621 - val_accuracy: 0.9223
Epoch 14/20
469/469 [=====] - 2s 4ms/step - loss: 0.2652 - accuracy: 0.9241 - val_loss: 0.2573 - val_accuracy: 0.9264
Epoch 15/20
469/469 [=====] - 2s 4ms/step - loss: 0.2592 - accuracy: 0.9234 - val_loss: 0.2511 - val_accuracy: 0.9277
Epoch 16/20
469/469 [=====] - 2s 4ms/step - loss: 0.2572 - accuracy: 0.9256 - val_loss: 0.2484 - val_accuracy: 0.9270
Epoch 17/20
469/469 [=====] - 2s 4ms/step - loss: 0.2529 - accuracy: 0.9277 - val_loss: 0.2462 - val_accuracy: 0.9284
Epoch 18/20
469/469 [=====] - 2s 4ms/step - loss: 0.2460 - accuracy: 0.9293 - val_loss: 0.2410 - val_accuracy: 0.9316
Epoch 19/20
469/469 [=====] - 2s 4ms/step - loss: 0.2394 - accuracy: 0.9329 - val_loss: 0.2280 - val_accuracy: 0.9349
Epoch 20/20
469/469 [=====] - 2s 4ms/step - loss: 0.2282 - accuracy: 0.9332 - val_loss: 0.2272 - val_accuracy: 0.9349
Test loss: 0.227160255098343
Test accuracy: 0.9348999857902527
```

```
In [10]: print(model.summary())

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
dense (Dense)                (None, 512)                 401920

dense_1 (Dense)              (None, 512)                 262656

dense_2 (Dense)              (None, 10)                  5130
=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

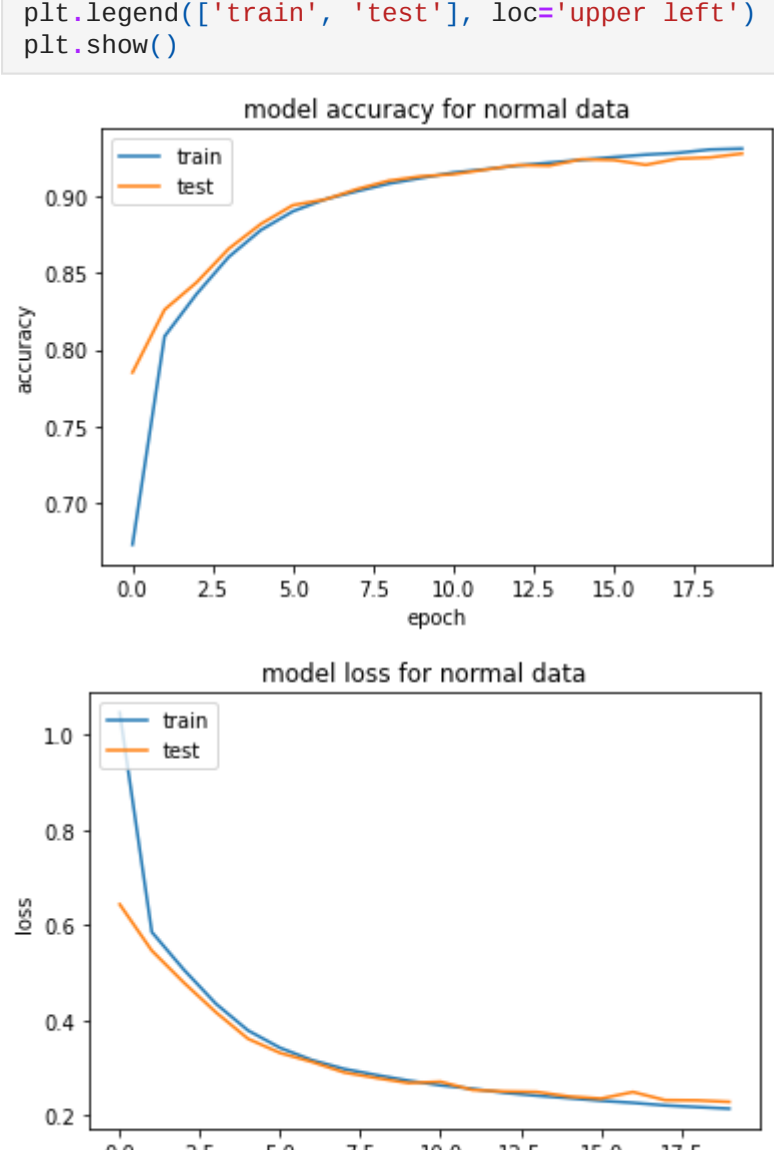
None
```

```
In [11]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

Test loss: 0.227160255098343
Test accuracy: 0.9348999857902527
```

```
In [22]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy for normal data')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss for normal data')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



Dimensionality Reduction using PCA

```
In [12]: x_train, y_train, x_valid, y_valid, x_test ,y_test = [], [], [], [], []
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train.reshape(60000, 28, 28)
x_test = x_test.reshape(10000, 28, 28)
```

```
In [13]: from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
x_train, x_valid, y_train, y_valid = train_test_split(x_train, y_train, test_size=0.2)

x_train = x_train.astype('float32')
x_valid = x_valid.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_valid /= 255
x_test /= 255

x_train = x_train.reshape(len(x_train), 28, 28)
x_test = x_test.reshape(10000, 28, 28)
y_train = keras.utils.to_categorical(y_train, 10)
y_valid = keras.utils.to_categorical(y_valid, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

```
In [14]: pca = PCA(n_components=10)
X_train_reduced = pca.fit_transform(X_train)
print("Explained variance = ", (100*(np.sum(pca.explained_variance_ratio_))))
print("The number of features are reduced to ", pca.n_components_)

Explained variance = 48.81797134876251
The number of features are reduced to 10
```

```
In [15]: X_test = x_test.reshape(10000, 784)
x_test_reduced = pca.transform(X_test)
X_valid = x_valid.reshape(12000, 784)
x_valid_reduced = pca.transform(X_valid)
```

```
In [16]: optim = keras.optimizers.SGD(lr=0.1)
```

```
In [17]: PCAModel = Sequential()
PCAModel.add(Dense(256, activation='sigmoid', kernel_initializer = 'he_normal'))
PCAModel.add(Dense(128, activation='sigmoid'))
PCAModel.add(Dense(128, activation='sigmoid'))
PCAModel.add(Dense(100, activation='sigmoid'))
PCAModel.add(Dense(100, activation='sigmoid'))
PCAModel.add(Dense(10, activation='sigmoid'))
optim = keras.optimizers.SGD(lr=0.1)
```

```
In [18]: PCAModel.compile(loss='categorical_crossentropy',
                      optimizer=keras.optimizers.SGD(lr=0.1, momentum=0.0, nesterov=False),
                      metrics=['accuracy'])
```

```
In [19]: history = PCAModel.fit(x_train_reduced, y_train, batch_size=32, epochs=20, verbose=2, validation_data=(x_valid_reduced, y_valid))

Epoch 1/20
1500/1500 - 4s - loss: 1.0462 - accuracy: 0.6733 - val_loss: 0.6432 - val_accuracy: 0.7853
Epoch 2/20
1500/1500 - 3s - loss: 0.5849 - accuracy: 0.8087 - val_loss: 0.5459 - val_accuracy: 0.8261
Epoch 3/20
1500/1500 - 3s - loss: 0.5058 - accuracy: 0.8363 - val_loss: 0.4786 - val_accuracy: 0.8438
Epoch 4/20
1500/1500 - 3s - loss: 0.4342 - accuracy: 0.8603 - val_loss: 0.4162 - val_accuracy: 0.8657
Epoch 5/20
1500/1500 - 3s - loss: 0.3781 - accuracy: 0.8778 - val_loss: 0.3606 - val_accuracy: 0.8817
Epoch 6/20
1500/1500 - 3s - loss: 0.3415 - accuracy: 0.8900 - val_loss: 0.3308 - val_accuracy: 0.8940
Epoch 7/20
1500/1500 - 3s - loss: 0.3156 - accuracy: 0.8978 - val_loss: 0.3118 - val_accuracy: 0.8977
Epoch 8/20
1500/1500 - 3s - loss: 0.2972 - accuracy: 0.9030 - val_loss: 0.2898 - val_accuracy: 0.9044
Epoch 9/20
1500/1500 - 3s - loss: 0.2844 - accuracy: 0.9080 - val_loss: 0.2781 - val_accuracy: 0.9099
Epoch 10/20
1500/1500 - 3s - loss: 0.2726 - accuracy: 0.9116 - val_loss: 0.2676 - val_accuracy: 0.9128
Epoch 11/20
1500/1500 - 3s - loss: 0.2627 - accuracy: 0.9152 - val_loss: 0.2703 - val_accuracy: 0.9141
Epoch 12/20
1500/1500 - 3s - loss: 0.2555 - accuracy: 0.9173 - val_loss: 0.2530 - val_accuracy: 0.9171
Epoch 13/20
1500/1500 - 3s - loss: 0.2470 - accuracy: 0.9197 - val_loss: 0.2499 - val_accuracy: 0.9200
Epoch 14/20
1500/1500 - 3s - loss: 0.2410 - accuracy: 0.9216 - val_loss: 0.2488 - val_accuracy: 0.9197
Epoch 15/20
1500/1500 - 3s - loss: 0.2358 - accuracy: 0.9234 - val_loss: 0.2393 - val_accuracy: 0.9238
Epoch 16/20
1500/1500 - 3s - loss: 0.2303 - accuracy: 0.9252 - val_loss: 0.2349 - val_accuracy: 0.9233
Epoch 17/20
1500/1500 - 3s - loss: 0.2260 - accuracy: 0.9268 - val_loss: 0.2487 - val_accuracy: 0.9202
Epoch 18/20
1500/1500 - 3s - loss: 0.2206 - accuracy: 0.9279 - val_loss: 0.2316 - val_accuracy: 0.9242
Epoch 19/20
1500/1500 - 3s - loss: 0.2171 - accuracy: 0.9301 - val_loss: 0.2311 - val_accuracy: 0.9250
Epoch 20/20
1500/1500 - 3s - loss: 0.2137 - accuracy: 0.9308 - val_loss: 0.2280 - val_accuracy: 0.9275
```

```
In [20]: print(model.summary())

Model: "sequential"

Layer (type)                Output Shape                Param #
=====
dense (Dense)                (None, 512)                 401920

dense_1 (Dense)              (None, 512)                 262656

dense_2 (Dense)              (None, 10)                  5130
=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

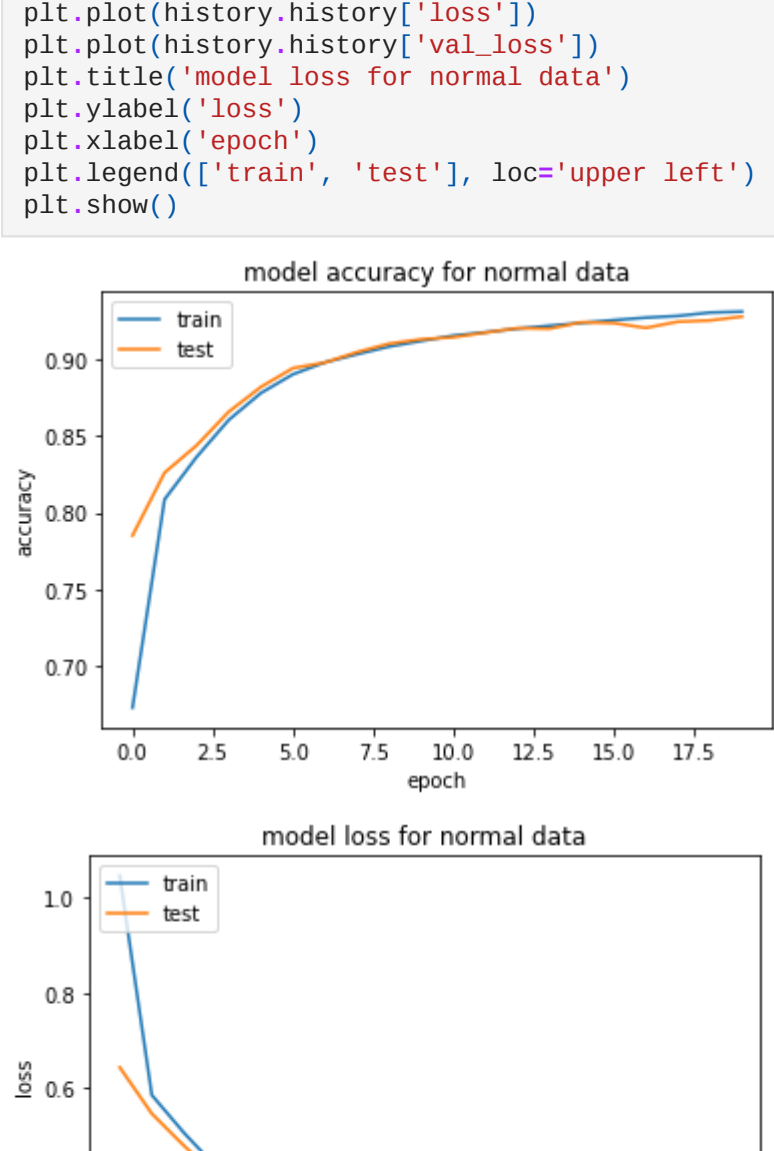
None
```

```
In [21]: score = PCAModel.evaluate(x_test_reduced, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

Test loss: 0.22815966606140137
Test accuracy: 0.9283000230789185
```

```
In [23]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy for normal data')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss for normal data')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [ ]:
```